

UCI Shopping EDA

Aaron Tang

2026-02-09

Link to dataset: <https://archive.ics.uci.edu/dataset/468/online+shoppers+purchasing+intention+dataset>

```
# Loading data
work_dir <- here()
raw_data <- read_csv(file.path(work_dir, "shopping", "data", "raw_data", "uci_shopping.csv"),
                      show_col_types = FALSE)
head(raw_data)

## # A tibble: 6 x 18
##   Administrative Administrative_Duration Informational Informational_Duration
##       <dbl>                 <dbl>          <dbl>                  <dbl>
## 1         0                  0             0                   0
## 2         0                  0             0                   0
## 3         0                  0             0                   0
## 4         0                  0             0                   0
## 5         0                  0             0                   0
## 6         0                  0             0                   0
## # i 14 more variables: ProductRelated <dbl>, ProductRelated_Duration <dbl>,
## #   BounceRates <dbl>, ExitRates <dbl>, PageValues <dbl>, SpecialDay <dbl>,
## #   Month <chr>, OperatingSystems <dbl>, Browser <dbl>, Region <dbl>,
## #   TrafficType <dbl>, VisitorType <chr>, Weekend <lgl>, Revenue <lgl>
```

Description of Variables

Continuous

- **Administrative:** Number of pages visited about account management
- **Administrative Duration:** Total amount of time (in seconds) spent on account management related pages
- **Informational:** Number of pages visited about website, communication and address information of the shopping site
- **Informational Duration:** Total amount of time spent on informational pages
- **Product Related:** Number of pages visited about product related pages
- **Product Related Duration:** Total amount of time spent on product related pages
- **Bounce Rate:** Average bounce rate value of the pages visited

- Formula: $BounceRate = 1 - EngagementRate$
- $EngagementRate$ = percentage of sessions that either lasted longer than 10 seconds, had key events, or two or more screen or page views
- Google Analytics assigns this value to every page of a website
- This feature takes the average of all the bounce rates value from all the pages the user visits in a single session
- **Exit Rate:** Average exit rate value of the pages visited
 - Formula: $ExitRate = \frac{Number\text{ of }Exits}{Number\text{ of }Page\text{ Views}} \times 100$
- **Page Value:** Average page value of the pages visited
 - Formula: $PageValue = \frac{TotalRevenue + TotalGoalValue}{UniquePageViews}$
 - $TotalRevenue$ = amount of money generated from a website page (i.e. transaction page)
 - $TotalGoalValue$ = value assigned to a specific page that is defined by the business
 - $UniquePageViews$ = number of unique user visits, only counted once per session
- **Special Day:** Closeness of the site visiting time to a special day

Categorical

- **OperatingSystems:** Operating system of the visitor
- **Browser:** Browser of the visitor
- **Region:** Geographic region from which the session has been started by the visitor
- **TrafficType:** Traffic source by which the visitor has arrived at the website
- **VisitorType:** Visitor type as “New Visitor”, “Returning Visitor”, and “Other”
- **Weekend:** Boolean value indicating whether the date of the visit is a weekend
- **Month:** Month value of the visit date
- **Revenue:** Class label indicating whether the visit has been finalized with a transaction

Summary

```
continuous_var <- c("Administrative", "Administrative_Duration", "Informational",
                     "Informational_Duration", "ProductRelated", "ProductRelated_Duration",
                     "BounceRates", "ExitRates", "PageValues")
```

```
categorical_var <- c("Month", "OperatingSystems", "Browser", "Region", "TrafficType",
                     "VisitorType", "Weekend", "Revenue")
```

```
shopping_uci <- raw_data %>%
  mutate(across(all_of(categorical_var), as.factor))
```

```
summary(shopping_uci)
```

```
##  Administrative  Administrative_Duration  Informational
##  Min.    : 0.000  Min.    : 0.00          Min.    : 0.0000
##  1st Qu.: 0.000  1st Qu.: 0.00          1st Qu.: 0.0000
##  Median  : 1.000  Median  : 7.50          Median  : 0.0000
##  Mean    : 2.315  Mean    : 80.82         Mean    : 0.5036
```

```

## 3rd Qu.: 4.000   3rd Qu.: 93.26          3rd Qu.: 0.0000
## Max.    :27.000   Max.    :3398.75        Max.    :24.0000
##
## Informational_Duration ProductRelated      ProductRelated_Duration
## Min.    : 0.00          Min.    : 0.00          Min.    : 0.0
## 1st Qu.: 0.00          1st Qu.: 7.00          1st Qu.: 184.1
## Median : 0.00          Median : 18.00         Median : 598.9
## Mean   : 34.47         Mean   : 31.73         Mean   : 1194.8
## 3rd Qu.: 0.00          3rd Qu.: 38.00         3rd Qu.: 1464.2
## Max.   :2549.38        Max.   :705.00         Max.   :63973.5
##
## BounceRates       ExitRates       PageValues      SpecialDay
## Min.   :0.0000000   Min.   :0.00000   Min.   : 0.000   Min.   :0.000000
## 1st Qu.:0.0000000  1st Qu.:0.01429  1st Qu.: 0.000  1st Qu.:0.00000
## Median :0.003112   Median :0.02516  Median : 0.000  Median :0.00000
## Mean   :0.022191   Mean   :0.04307  Mean   : 5.889  Mean   :0.06143
## 3rd Qu.:0.016813   3rd Qu.:0.05000  3rd Qu.: 0.000  3rd Qu.:0.00000
## Max.   :0.200000   Max.   :0.20000  Max.   :361.764  Max.   :1.00000
##
## Month          OperatingSystems     Browser        Region        TrafficType
## May           :3364   2       :6601   2       :7961   1       :4780   2       :3913
## Nov           :2998   1       :2585   1       :2462   3       :2403   1       :2451
## Mar           :1907   3       :2555   4       :736    4       :1182   3       :2052
## Dec           :1727   4       : 478   5       :467    2       :1136   4       :1069
## Oct           : 549   8       : 79    6       :174    6       : 805   13      : 738
## Sep           : 448   6       : 19    10      :163    7       : 761   10      : 450
## (Other):1337 (Other): 13      (Other):367 (Other):1263 (Other):1657
##
## VisitorType      Weekend       Revenue
## New_Visitor     : 1694 FALSE:9462  FALSE:10422
## Other          :  85  TRUE :2868   TRUE : 1908
## Returning_Visitor:10551
##
##
##
```


Bar Charts/Histograms of Variables

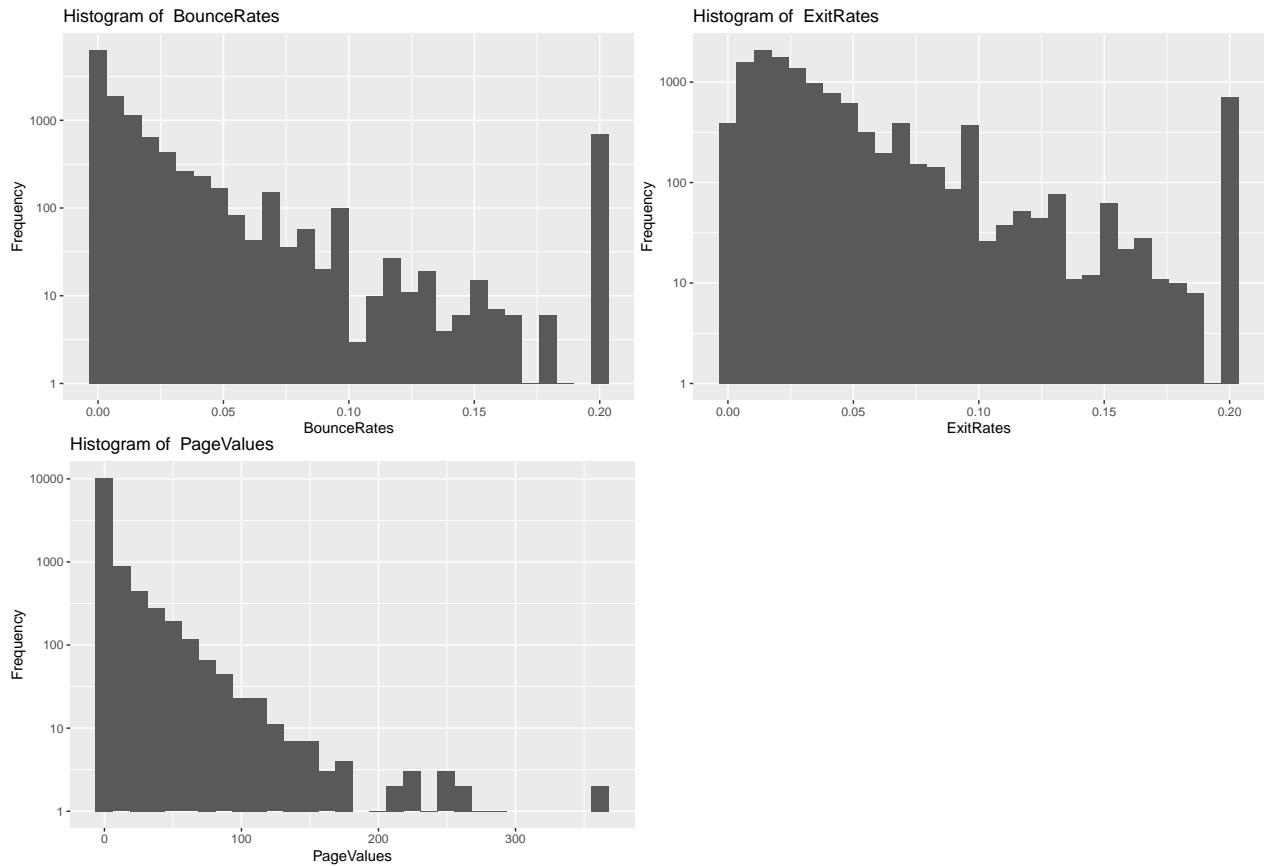
```

# Histograms of Continuous Variables
par(mfrow = c(2, 3))

for (i in 1:length(continuous_var)) {
  current = shopping_uci[[continuous_var[i]]]
  plot <- ggplot(data = shopping_uci, mapping = aes(x = current)) +
    geom_histogram() +
    labs(
      y = "Frequency",
      x = continuous_var[i],
      title = paste("Histogram of ", continuous_var[i])
    ) +
    scale_y_log10()
}
```

```
    print(plot)
}
```







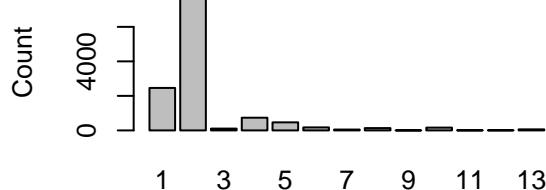
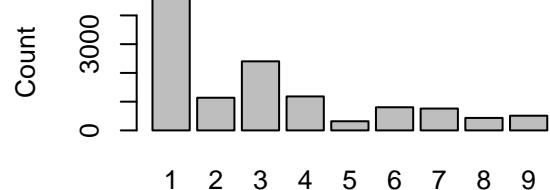

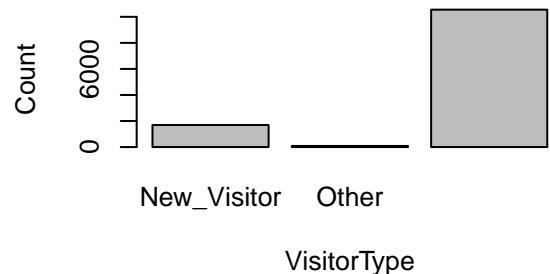


Notes:

- I logged the scale of the y-axes because they were all heavily zero-inflated
- After log transforming the scale, it appears the duration graphs follow a somewhat exponential distribution
- The rates appear to be a little left skewed

```
# Bar Charts of Categorical Variables
par(mfrow = c(2, 2))

for (i in 1:length(categorical_var)) {
  current_var_count <- shopping_uci %>%
    group_by(across(all_of(categorical_var[i]))) %>%
    summarize(count = n())

  barplot(height = current_var_count$count,
          names.arg = current_var_count %>% pull(1),
          xlab = categorical_var[i],
          ylab = "Count",
          main = paste("Frequency of", categorical_var[i]))
}
```

Frequency of Month**Frequency of OperatingSystems****Frequency of Browser****Frequency of Region****Frequency of TrafficType****Frequency of VisitorType****Frequency of Weekend****Frequency of Revenue**

Seeing if combining durations is useful


```
# Creating new dataset with total durations
shopping_uci_total_durations <- shopping_uci %>%
  mutate(Total_Duration = Administrative_Duration + Informational_Duration + ProductRelated_Duration)

shopping_uci_total_durations

## # A tibble: 12,330 x 19
##   Administrative_Duration Informational_Duration ProductRelated_Duration
##   <dbl>                 <dbl>                  <dbl>
## 1 0                     0                      0
## 2 0                     0                      0
## 3 0                     0                      0
## 4 0                     0                      0
## 5 0                     0                      0
## 6 0                     0                      0
## 7 0                     0                      0
## 8 1                     0                      0
## 9 0                     0                      0
## 10 0                    0                      0
## # i 12,320 more rows
## # i 15 more variables: ProductRelated_Duration <dbl>, ProductRelated_Duration <dbl>,
## #   BounceRates <dbl>, ExitRates <dbl>, PageValues <dbl>, SpecialDay <dbl>,
## #   Month <fct>, OperatingSystems <fct>, Browser <fct>, Region <fct>,
## #   TrafficType <fct>, VisitorType <fct>, Weekend <fct>, Revenue <fct>,
## #   Total_Duration <dbl>

# Checking to see if total duration correlates with other individual durations
corr_values <- shopping_uci_total_durations %>%
  select(c(Administrative_Duration, Informational_Duration, ProductRelated_Duration, Total_Duration)) %>%
  cor()

corrplot(
  corr_values,
  diag = FALSE,
  method = "color",
  tl.col = "black"
)
```




- Total_Duration seems to correlate with the other individual durations, especially ProductRelated_Duration
- From this, it may be safe to sum up the individual durations into one variable, as it looks like we would not lose that much information if we perform this operation

Multivariate analysis


```
# Box plots
total_duration <- ggplot(data = shopping_uci_total_durations, mapping = aes(x = Revenue, y = Total_Duration)) +
  geom_boxplot()
bounce <- ggplot(data = shopping_uci_total_durations, mapping = aes(x = Revenue, y = BounceRates)) +
  geom_boxplot()
exit <- ggplot(data = shopping_uci_total_durations, mapping = aes(x = Revenue, y = ExitRates)) +
  geom_boxplot()
page <- ggplot(data = shopping_uci_total_durations, mapping = aes(x = Revenue, y = PageValues)) +
  geom_boxplot()

grid.arrange(total_duration, bounce, exit, page, ncol = 2)
```




Nothing really useful here.


```
# Scatter plots
ggplot(data = shopping_uci_total_durations, mapping = aes(x = Total_Duration, y = BounceRates,
                                                       color = Revenue)) +
  geom_point(alpha = 0.5) +
  coord_cartesian(xlim = c(0, 1000))
```



```
ggplot(data = shopping_uci_total_durations, mapping = aes(x = Total_Duration, y = ExitRates,
                                                       color = Revenue)) +
  geom_point(alpha = 0.5) +
  coord_cartesian(xlim = c(0, 1000))
```



```
ggplot(data = shopping_uci_total_durations, mapping = aes(x = Total_Duration, y = PageValues,
                                                       color = Revenue)) +
  geom_point(alpha = 0.5) +
  coord_cartesian(xlim = c(0, 1000))
```



Nothing really useful can be extracted from this as well

```
# Correlation plot between rates
corr_values <- shopping_uci_total_durations %>%
  select(c(BounceRates, ExitRates, PageValues)) %>%
  cor()

corrplot(
  corr_values,
  diag = FALSE,
  method = "color",
  tl.col = "black"
)
```

